WORKFLOW PORTFOLIO

26 workflows, one habit.

Mathieu Tellene · automation and AI built at Cabify

Every n8n workflow I have published and left running in production, one per page, captured from the live canvas. Five project areas, 26 live workflows, 537 nodes. The numbers on each page describe the flow itself — how many nodes, how much of it is AI, where it branches — not the business results, which belong in a different conversation.

26 workflows live in production	**537** nodes across all of them	**41** AI & LLM nodes	**34** independent triggers	**17** sub-workflow calls

Systems wired in: Claude (LLM gateway) · Google Docs · Google Drive · Google Sheets · HTTP / REST API · IMAP · Image processing · PDF parsing · Redis · Slack · Tableau · WhatsApp · n8n DataTable

What is inside

PROJECT

WhatsApp Booking Agent

WhatsApp Booking Agent · n8n

One conversational agent and the satellite flows around it: address resolution against 44 city hubs, quoting, booking, cancellation, live journey tracking, the static map, the invoice and the post-trip support ticket. This is the project that produced the company's first fully automated ride.

10	**183**	**8**	**11**	**17**	**14**
Workflows	Nodes	AI nodes	Triggers	Branch points	Sub-workflow calls

Systems wired in: Claude (LLM gateway) · HTTP / REST API · IMAP · Redis · Slack · WhatsApp · n8n DataTable

WhatsApp booking agent · main conversation

The conversation itself. A WhatsApp trigger routes every inbound message and one AI agent, backed by Redis memory, decides which of five sub-workflows to call: resolve the address, quote the trip, book it, cancel it, or hand over to support. Pickup and product lists are rebuilt from cached data and sent as interactive buttons, so the rider picks instead of typing.

54	8	1	7	3	7
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

Wired to: n8n DataTable · WhatsApp · HTTP / REST API · Redis · Claude (LLM gateway)

> **Why it is built this way.** The agent owns the conversation but never calls the booking API itself. Every action is a separate sub-workflow, so a failing endpoint cannot take the conversation down with it.

Address and hub resolution

Turns whatever the rider typed or shared into coordinates the booking API will accept. It reverse-geocodes a shared location, runs a place search on a written address, matches the result against the city hub list and caches the resolved point in Redis.

15 Nodes	**0** AI nodes	**1** Triggers	**2** Branch points	**6** Code / mapping	**0** Sub-workflow calls

Wired to: HTTP / REST API · Redis · n8n DataTable

> **Why it is built this way.** Ambiguity is handed back to the agent as a question instead of being resolved by guessing. A wrong pickup point costs a cancelled ride.

Trip quote

Prices the trip. Reads the resolved origin and destination, asks the API for available products and fares, and returns a clean list to the agent while caching the chosen option for the booking step.

10	0	1	0	5	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

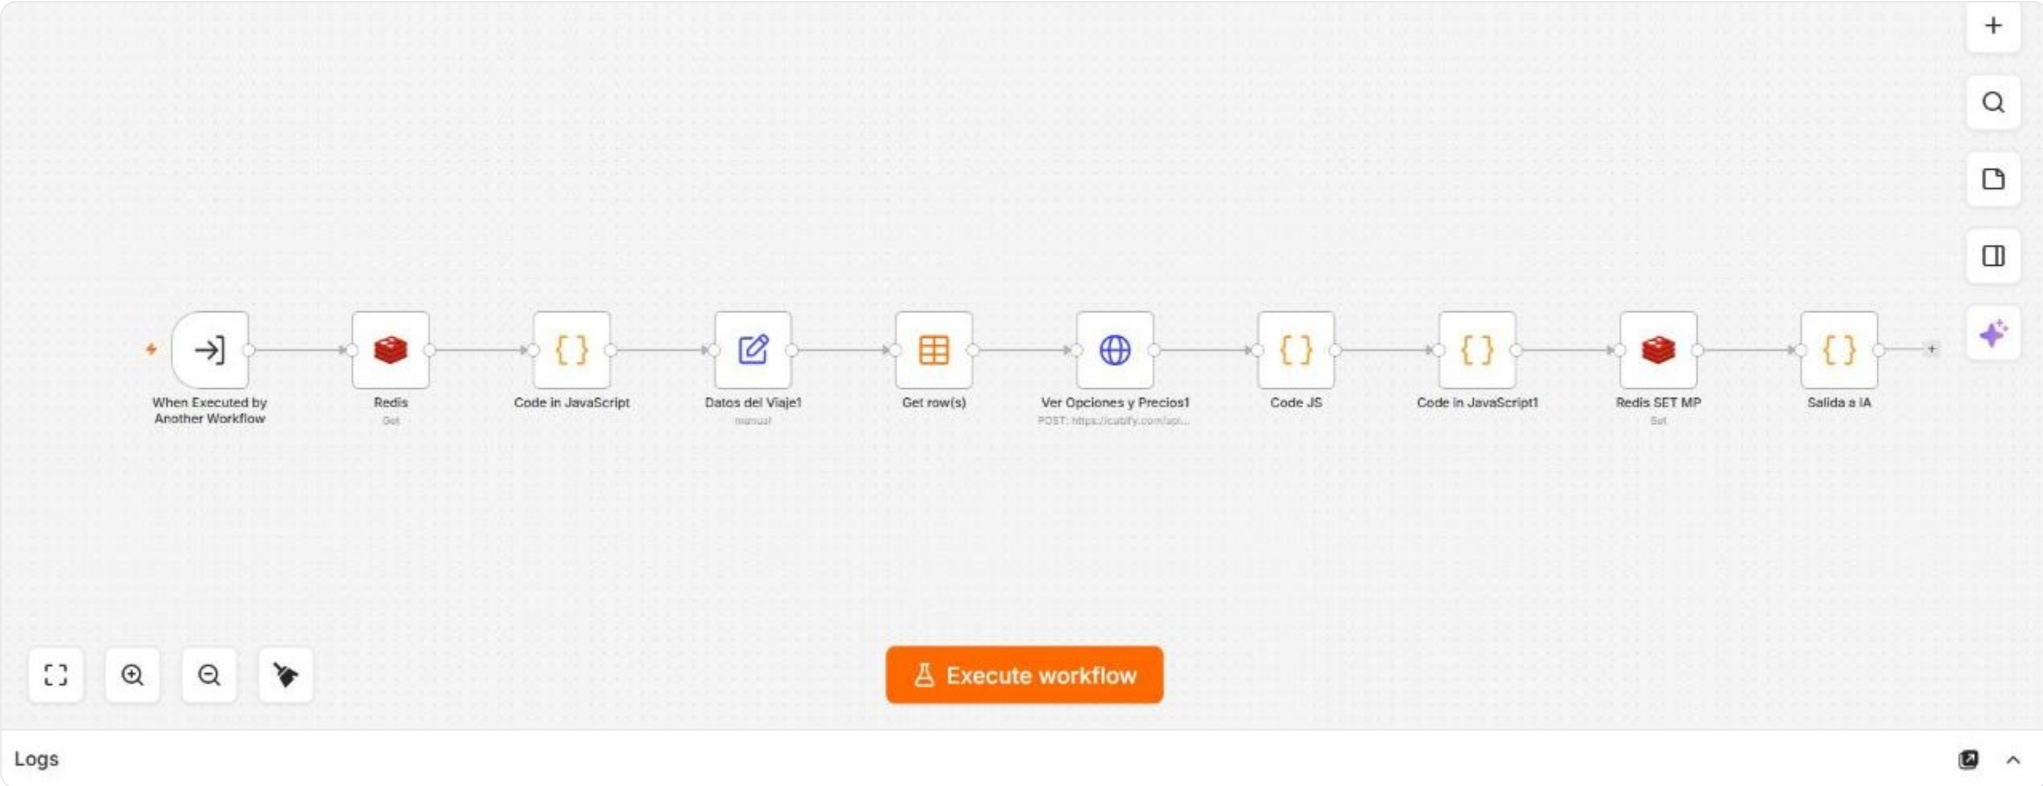

Wired to: Redis · HTTP / REST API · n8n DataTable

Trip creation

Books the ride with the cached quote, waits for the driver assignment and returns the share link the rider uses to follow the car, writing the trip into the shared table on the way out.

11	**0**	**1**	**0**	**3**	**0**
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

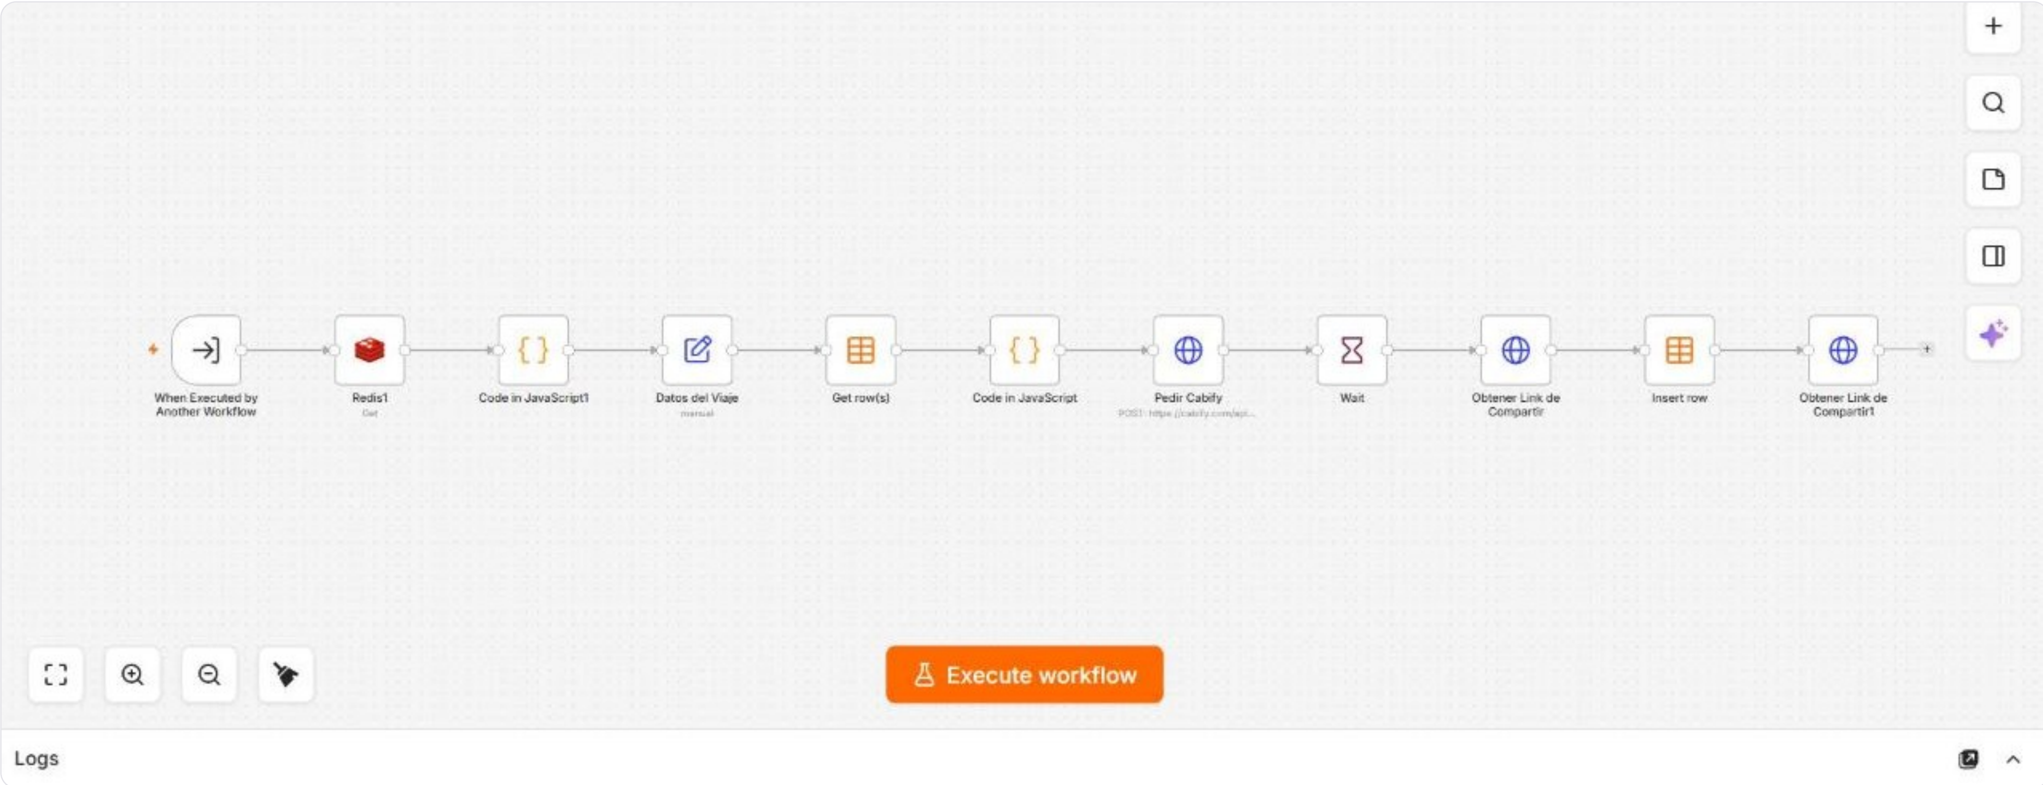

Wired to: HTTP / REST API · n8n DataTable · Redis

Trip cancellation

Cancels a booked ride: authenticates, sends the cancellation and verifies it actually went through before reporting success back to the agent.

6	0	1	0	0	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

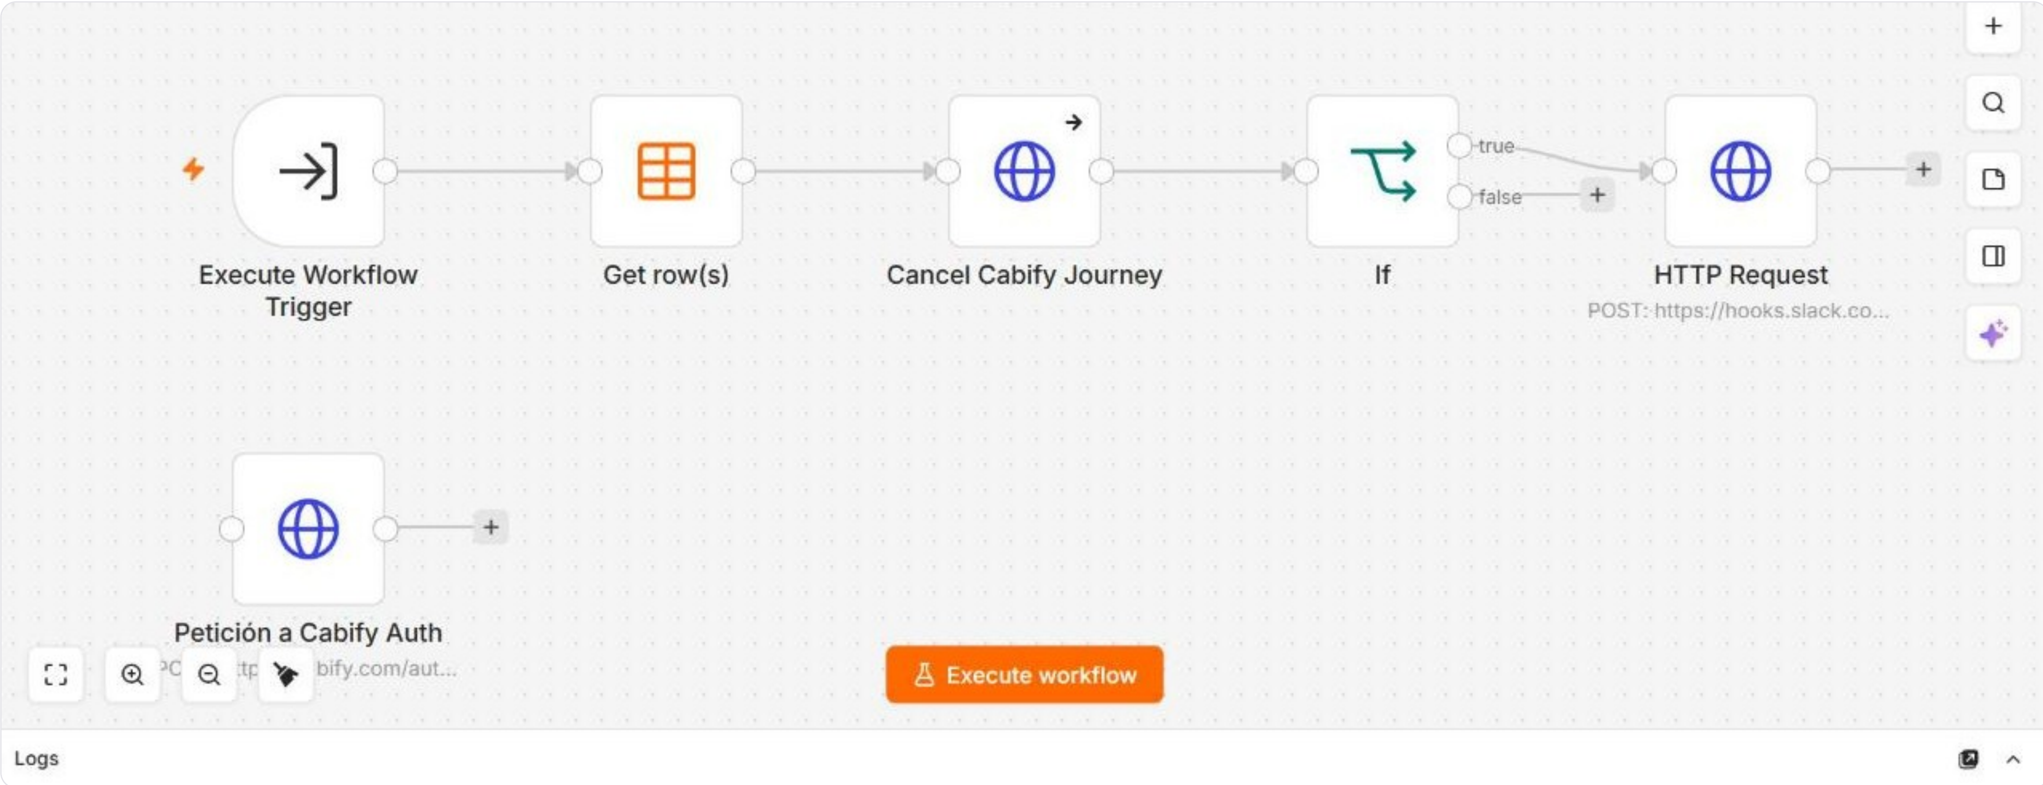

Wired to: HTTP / REST API · n8n DataTable

Journey tracking and post-trip

Everything that happens after a car is assigned. It polls the journey state, sends the rider the right message at each stage, records every state change in the trip tables, opens a support ticket when something goes wrong and posts to Slack when a new booking is created.

52	**0**	**1**	**5**	**0**	**6**
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

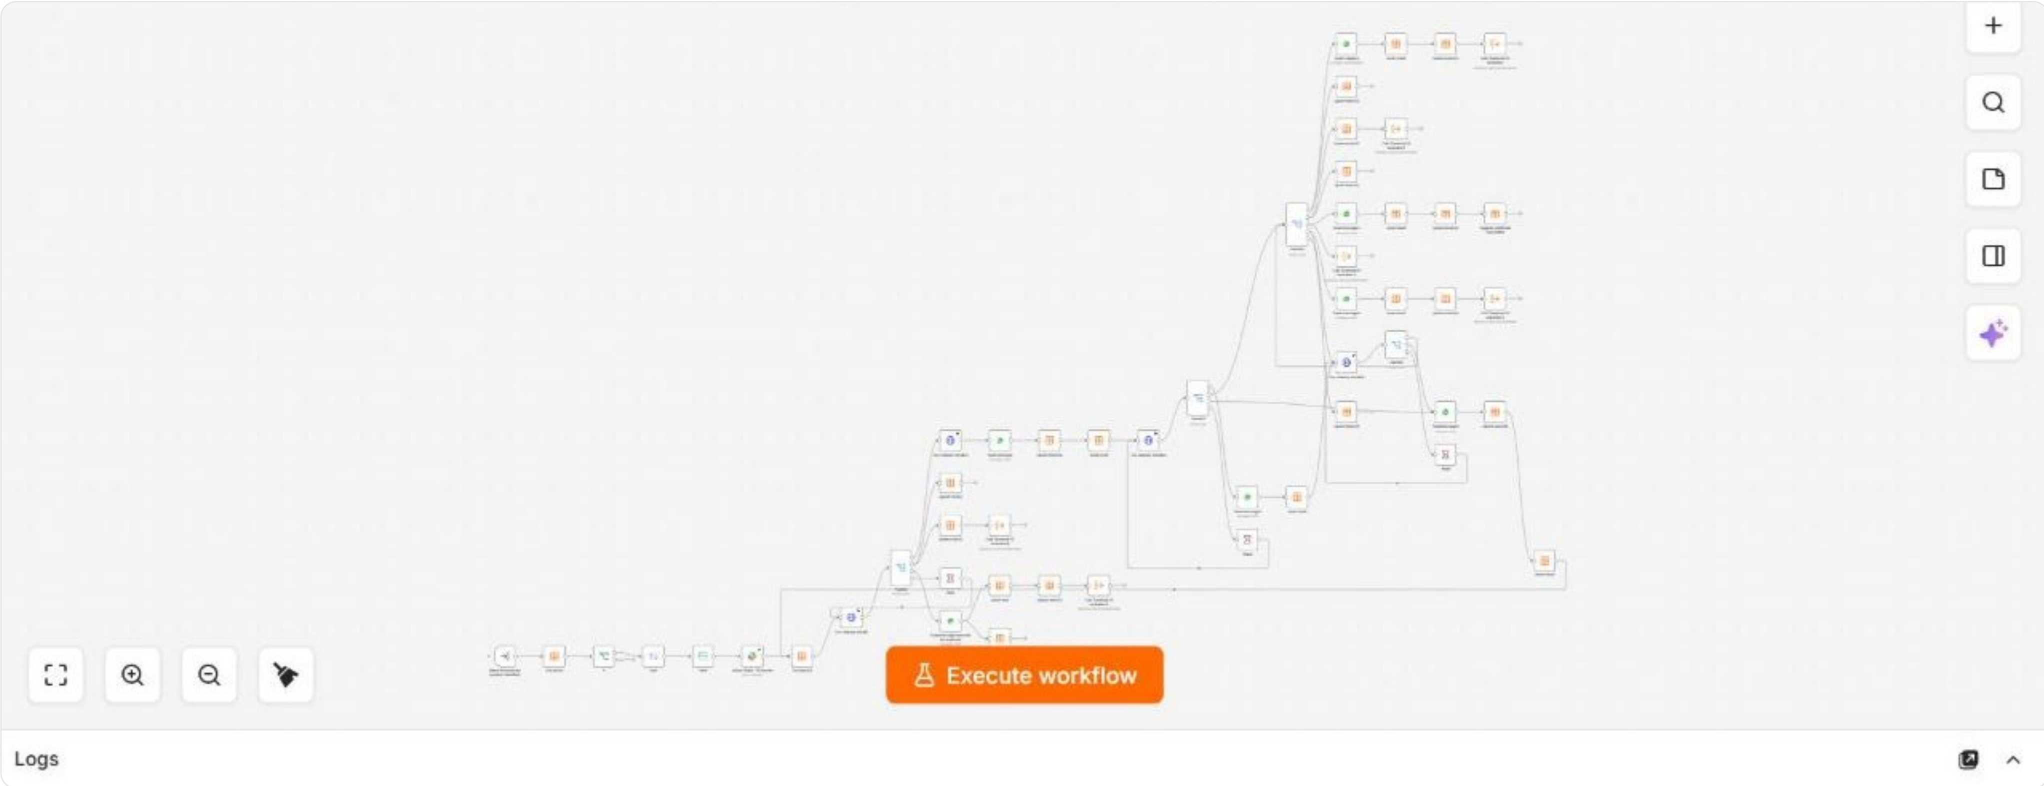

Wired to: n8n DataTable · WhatsApp · HTTP / REST API · Slack

> **Why it is built this way.** State lives in the data tables, not in the flow. If the flow restarts mid-trip the next poll picks up exactly where it left off.

API token refresh

A three-node scheduled job that keeps the API token fresh in the shared table, so none of the booking flows ever has to stop and authenticate in the middle of a conversation.

3 Nodes	**0** AI nodes	**1** Triggers	**0** Branch points	**0** Code / mapping	**0** Sub-workflow calls

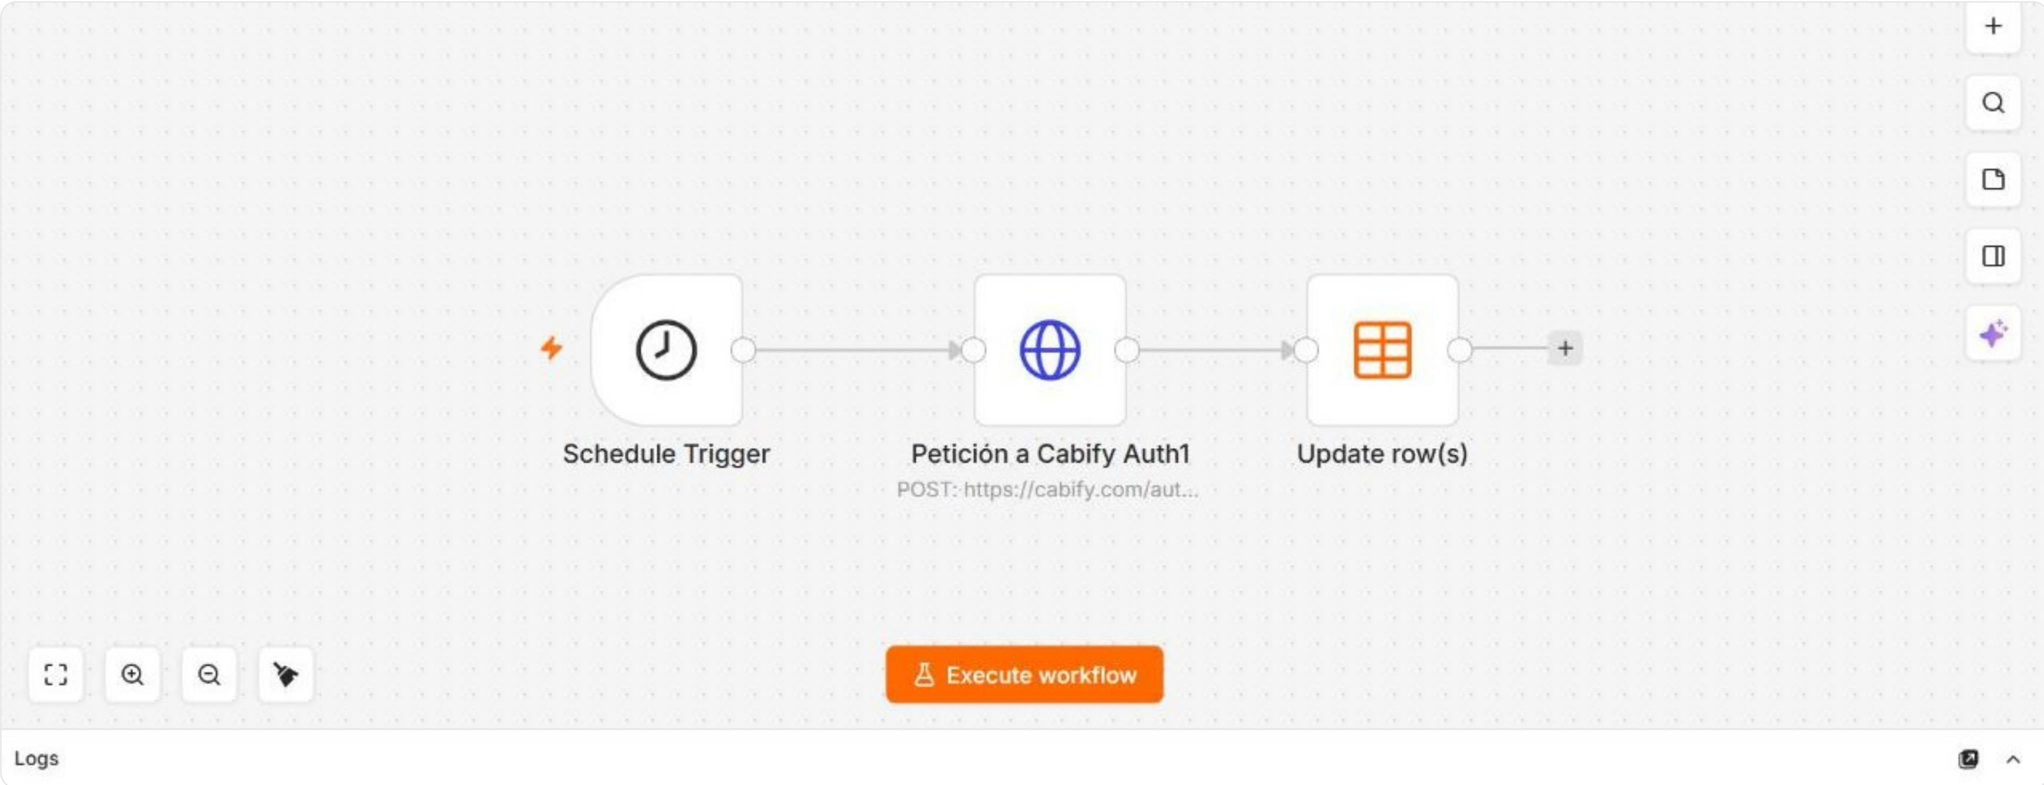

Wired to: HTTP / REST API · n8n DataTable

Static map service

A webhook that reads the cached pickup point and returns a static map image, so the rider sees where the car is coming to instead of reading a pair of coordinates.

5	0	2	0	1	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

Wired to: Redis · HTTP / REST API

Support ticket bridge

Assembles the support ticket: pulls the rider's conversation and trip rows, formats them and hands the result to the ticket-creation flow, so an escalation arrives with the full context already attached.

6 Nodes	0 AI nodes	1 Triggers	1 Branch points	1 Code / mapping	1 Sub-workflow calls

Wired to: n8n DataTable

Invoice delivery

Watches an inbox for the invoice, extracts the PDF, matches it to the right trip by reference and delivers it back to the rider on WhatsApp.

21	0	1	2	7	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

Wired to: n8n DataTable · WhatsApp · HTTP / REST API · IMAP

cabify

PROJECT

AQM — Asset Quality Management

AQM · Asset Quality Management · n8n

The whole quality loop for the ESP fleet. Capture services behind four live forms, the vision analysis engine that scores every photo against the official reference, the daily flows that keep fleet data and licence validation current, and the single database that turns an audit into a bonus payment.

10	143	14	15	15	1
Workflows	Nodes	AI nodes	Triggers	Branch points	Sub-workflow calls

Systems wired in: Claude (LLM gateway) · Google Drive · Google Sheets · HTTP / REST API · Image processing · Slack · n8n DataTable

Vehicle livery analysis engine

The audit engine. For each vehicle it downloads the submitted photos, prepares every image and runs a vision model per side. Where an official reference exists it downloads it, builds a real-versus-reference collage at matching size and asks a strict assessor for a verdict; where it does not, it falls back to a subjective audit. Seven branches merge into one consolidated result per vehicle.

81 Nodes	**10** AI nodes	**2** Triggers	**10** Branch points	**18** Code / mapping	**1** Sub-workflow calls

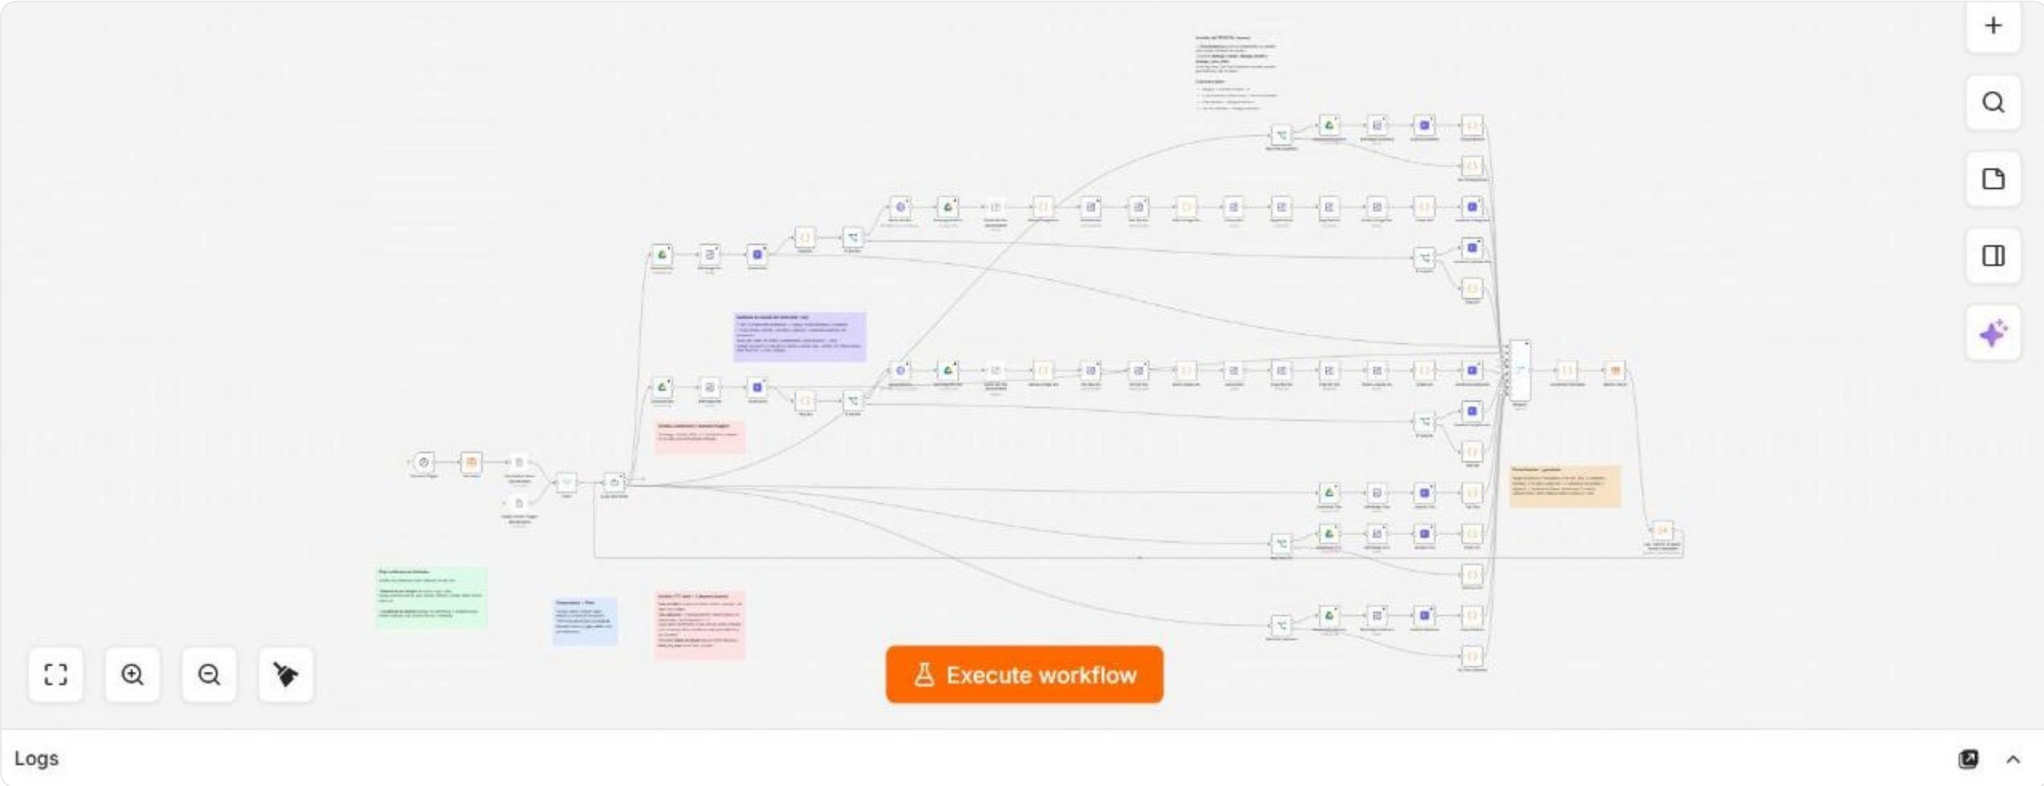

Wired to: Image processing · Google Drive · n8n DataTable · HTTP / REST API · Google Sheets · Claude (LLM gateway)

> **Why it is built this way.** One model call per image instead of five, and the collage is assembled inside the flow so the model always compares like with like. That is where most of the 80% cost reduction came from.

Plate OCR service

A webhook the capture form calls while the driver is still on screen: it reads the plate off the photo and returns it, so the form can flag a mismatch immediately.

5	1	2	0	2	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

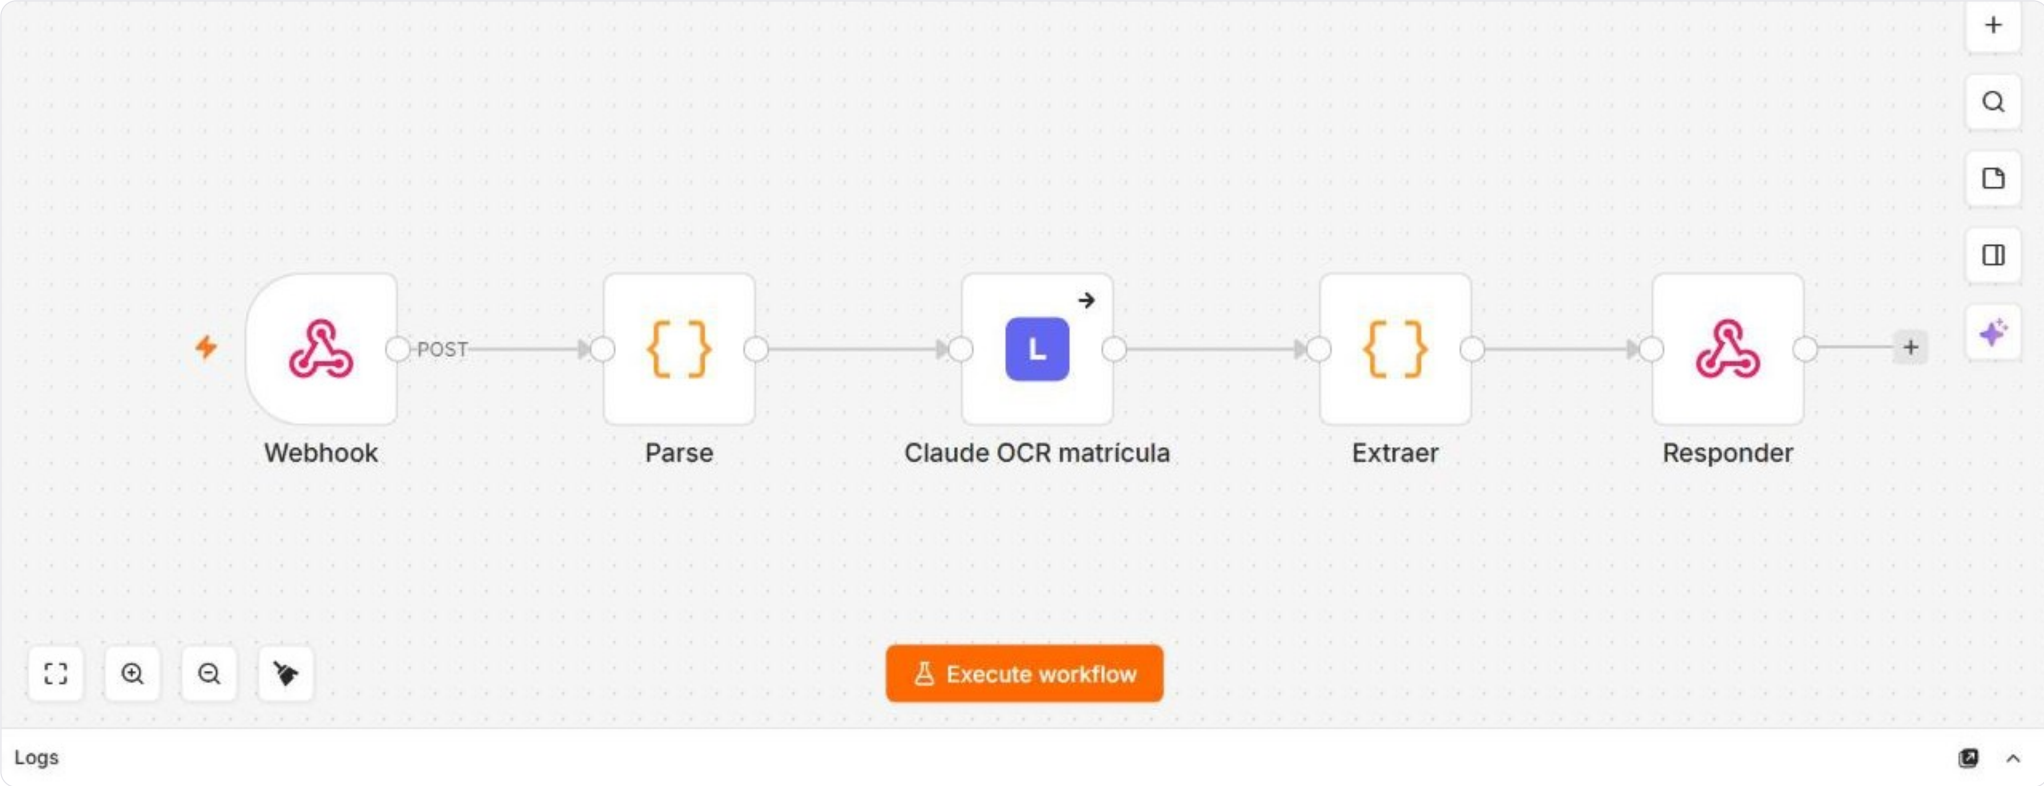

Wired to: Claude (LLM gateway)

Capture submission

What happens when a driver presses send: the files go to Drive, a row is written to the sheet and the data table, the plate is read again from the photo, and Slack is notified when the typed plate and the photographed plate disagree.

15 Nodes	**1** AI nodes	**2** Triggers	**1** Branch points	**6** Code / mapping	**0** Sub-workflow calls

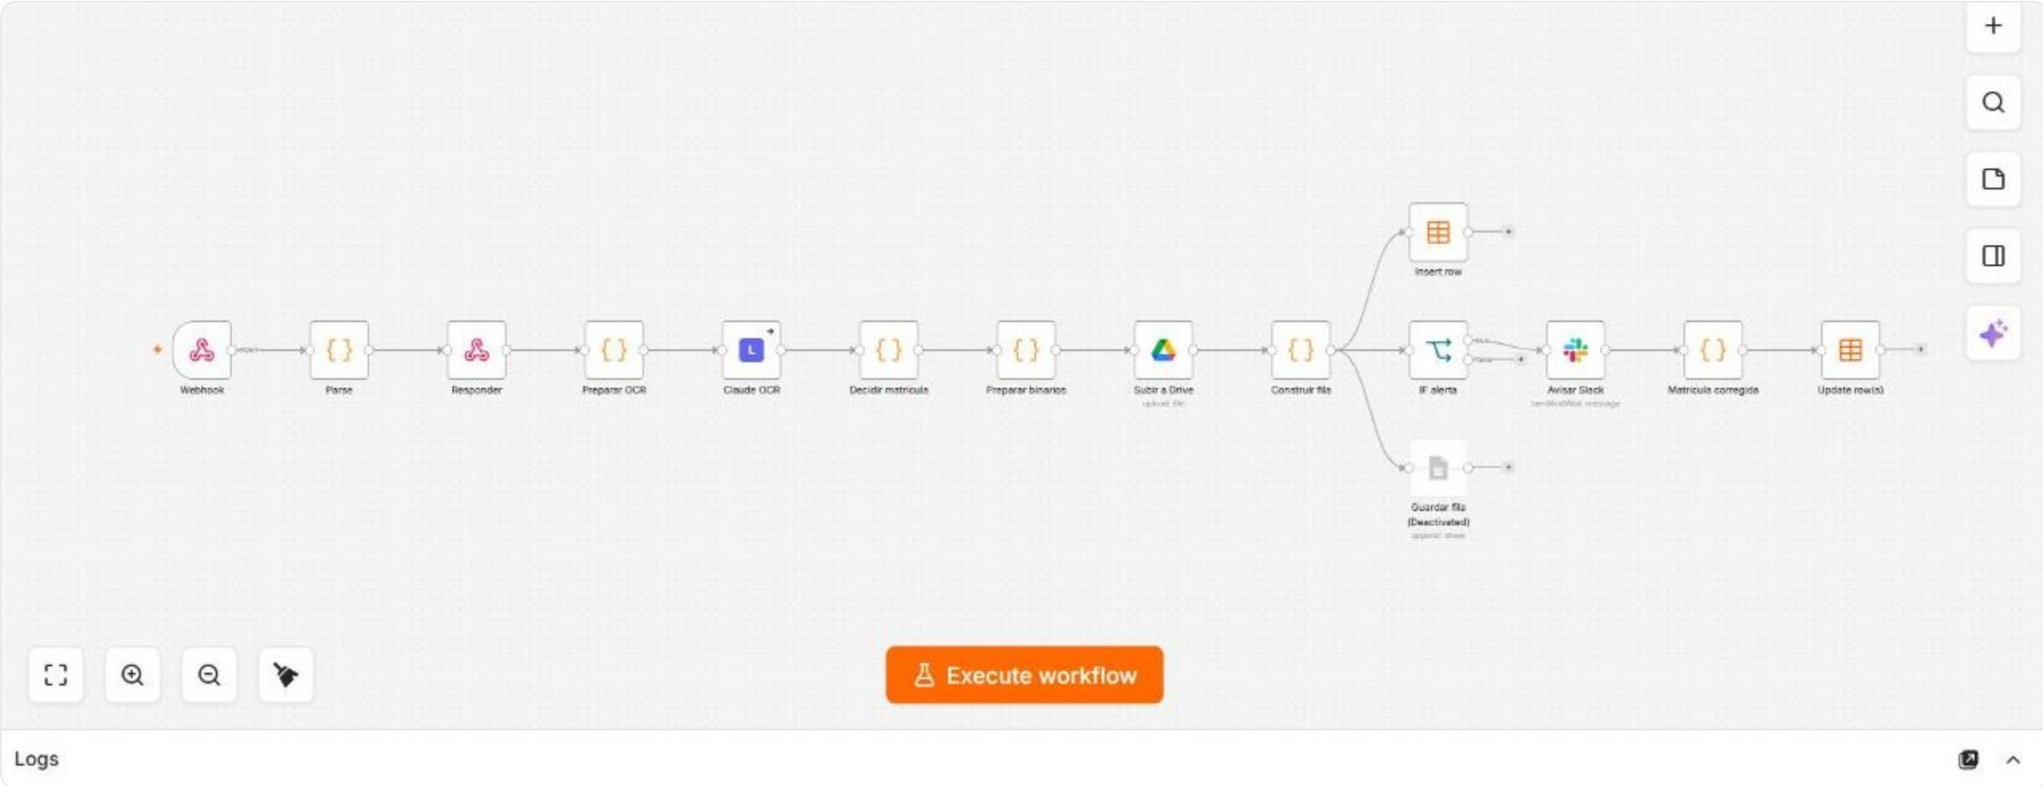

Wired to: n8n DataTable · Google Sheets · Google Drive · Slack · Claude (LLM gateway)

> **Why it is built this way.** The OCR warns, it never blocks. A false negative here would leave a driver unpaid, so the final call stays with a human.

Plate autocomplete

Serves the live plate list to the capture form, normalised, so a driver can only submit a plate that actually exists in the fleet.

4	0	2	0	1	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

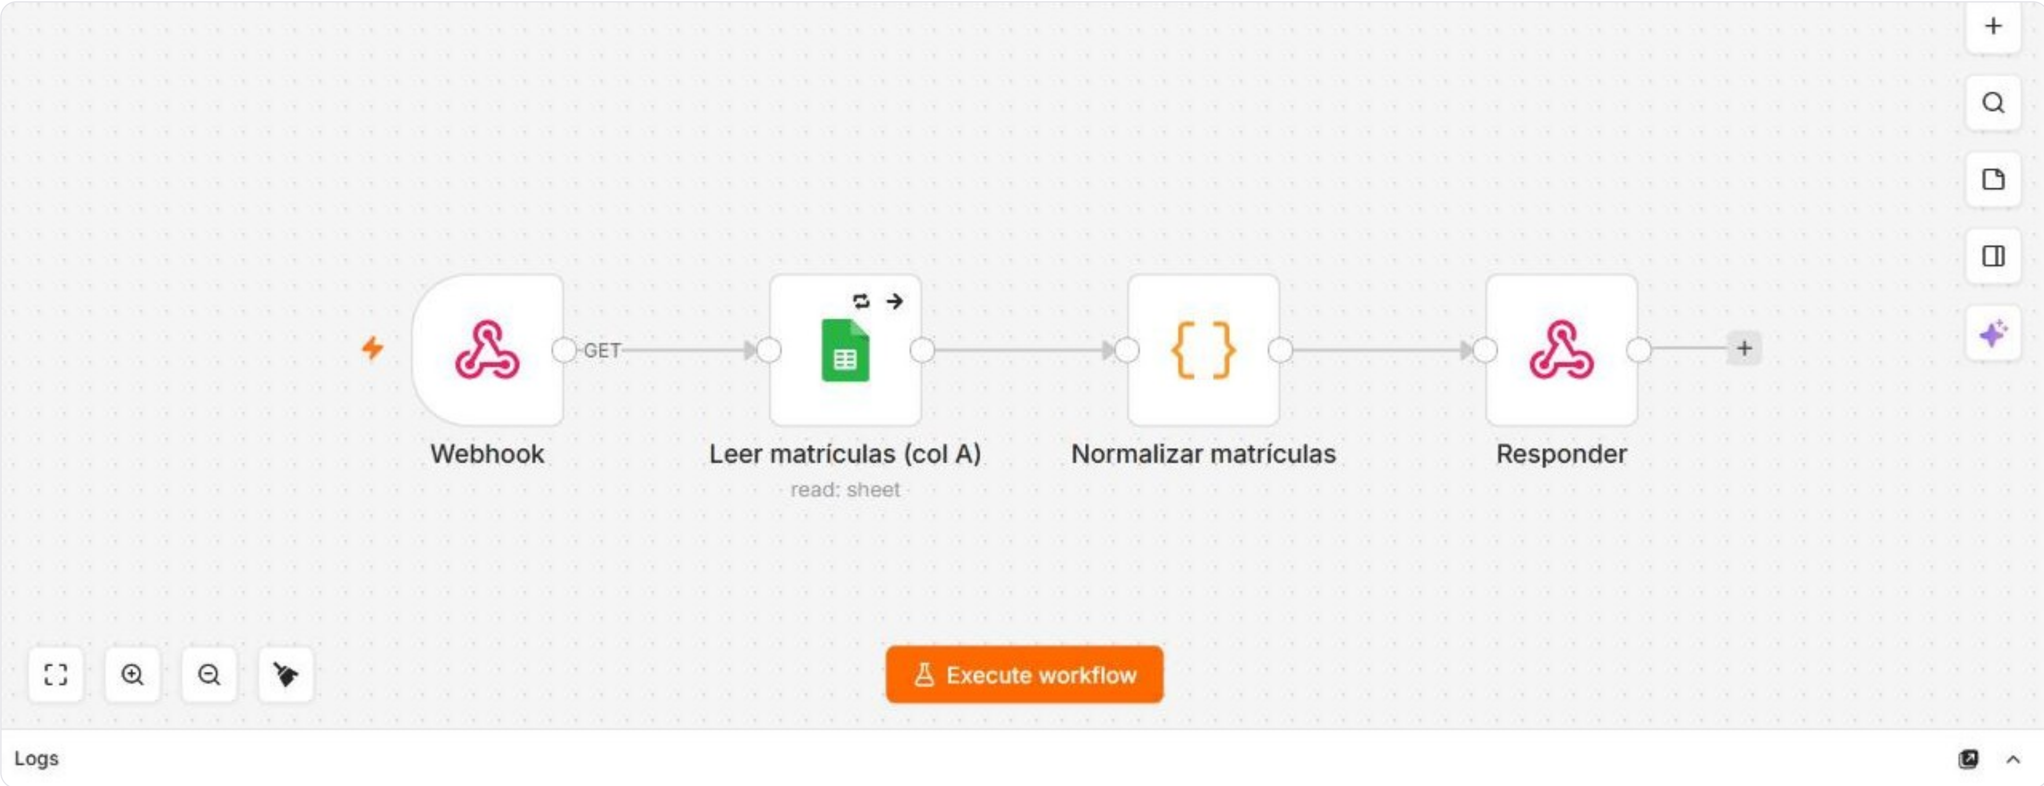

Wired to: Google Sheets

Monthly activity refresh

Picks up the monthly activity export the moment it lands in Drive, deduplicates by plate keeping the most recent record, and updates the master sheet.

9 Nodes

0 AI nodes

2 Triggers

1 Branch points

0 Code / mapping

0 Sub-workflow calls

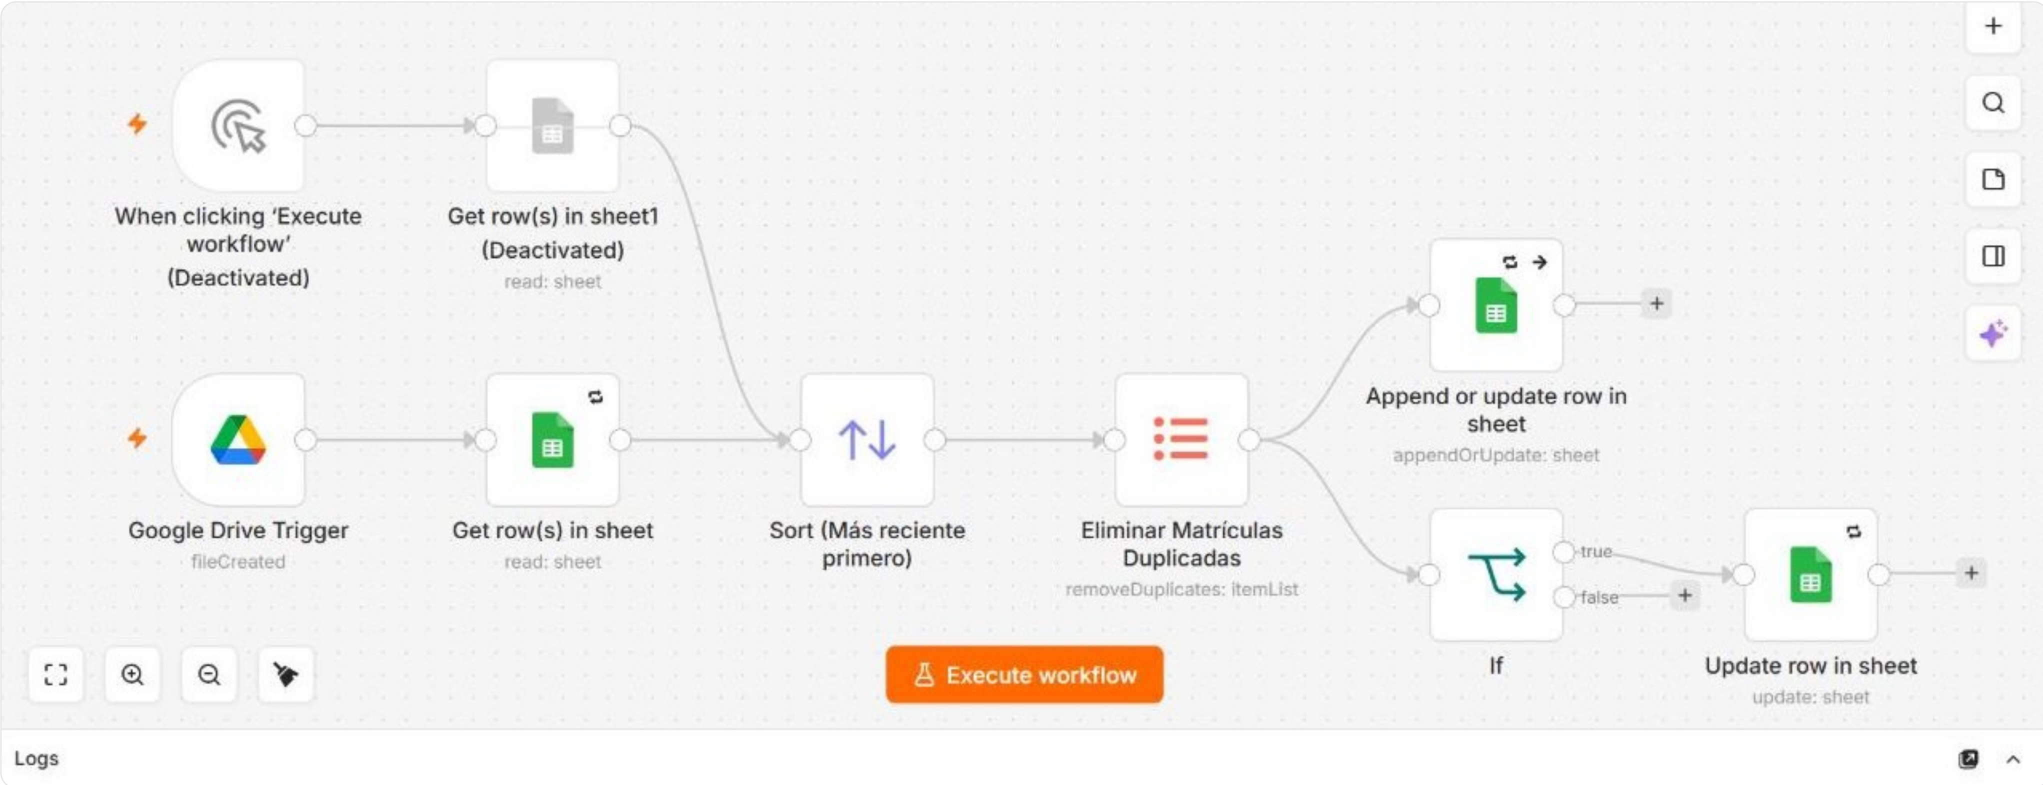

Wired to: Google Sheets

Licence certificate sync

Drive-triggered sync that keeps each vehicle's transport-licence certificate current in the master sheet, filtering out the rows that have not changed.

4 Nodes	**0** AI nodes	**1** Triggers	**0** Branch points	**0** Code / mapping	**0** Sub-workflow calls

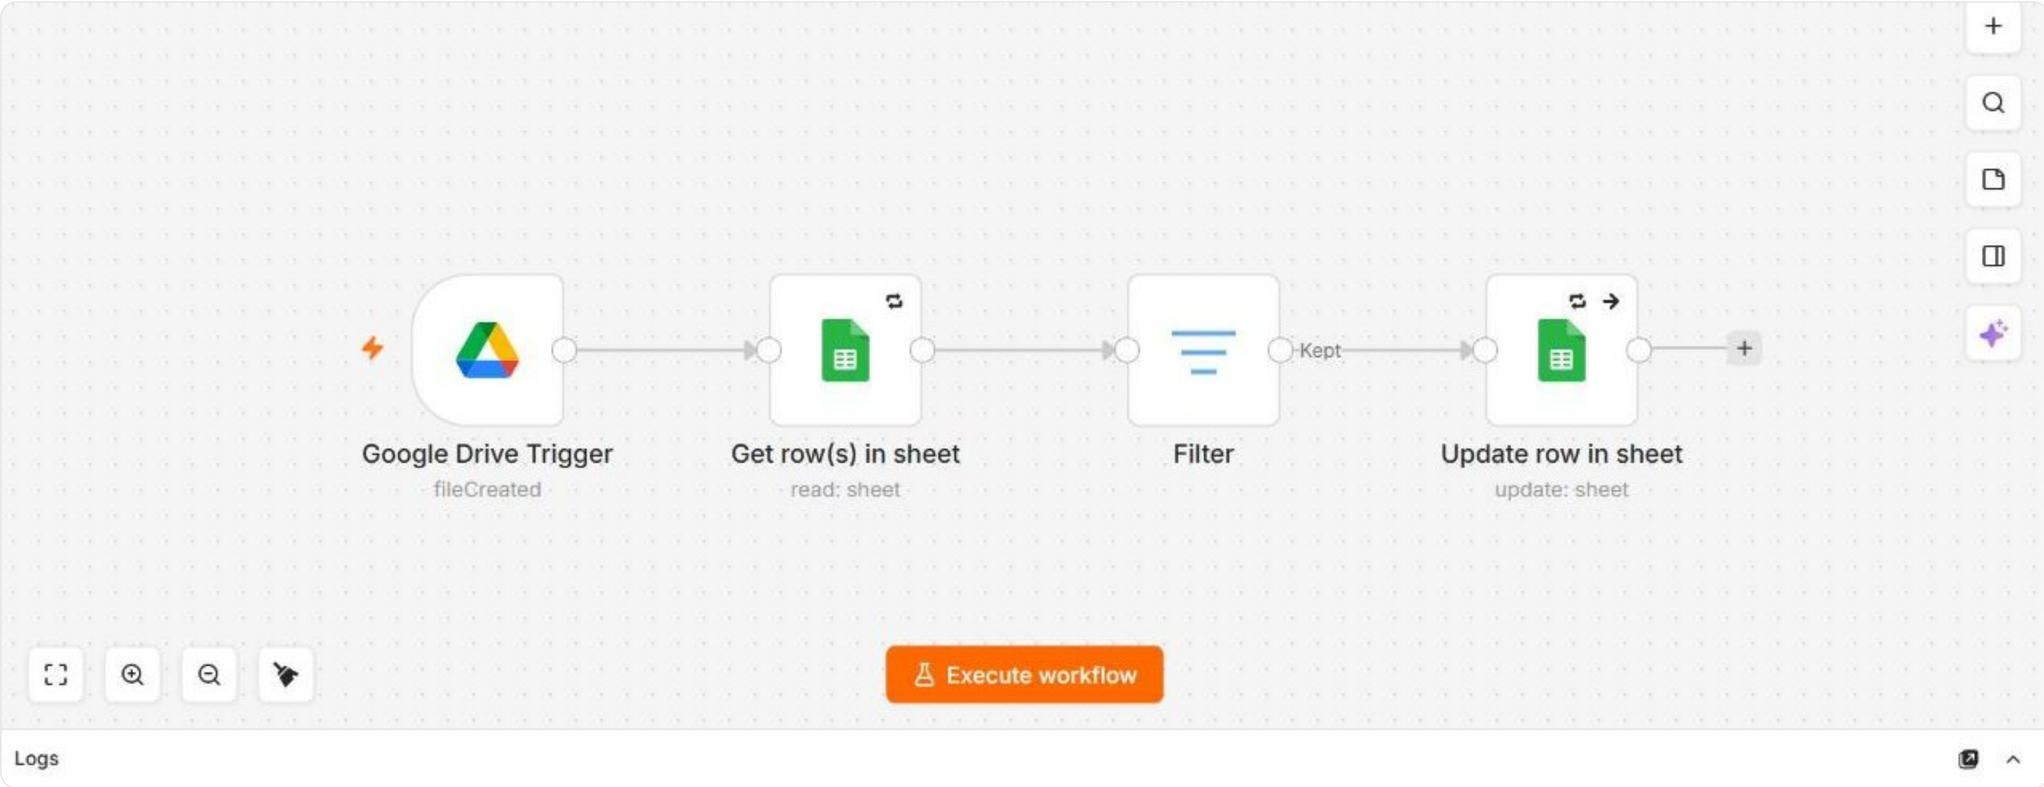

Wired to: Google Sheets

Licence validation sync

The companion sync that refreshes licence validation status, so an audit never approves a vehicle whose licence has lapsed.

3	0	1	0	0	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

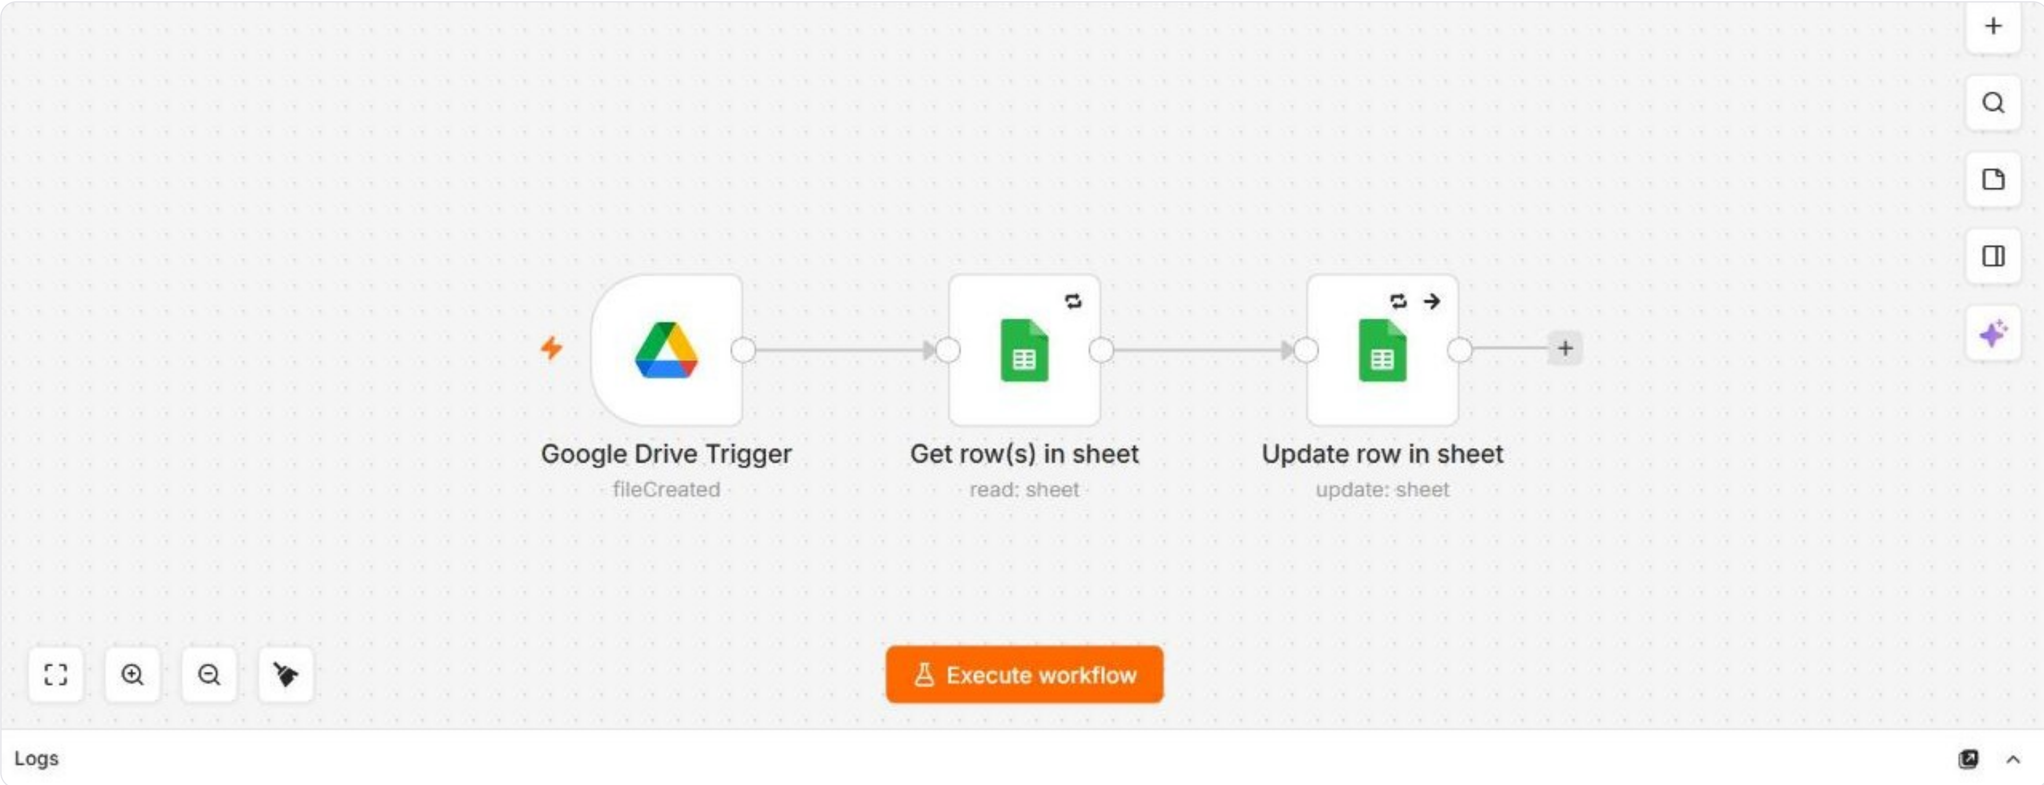

Wired to: Google Sheets

Fleet partner export intake

A webhook that receives the fleet partner's daily export, extracts the file and writes it straight into the fleet sheet. Three nodes, no manual step.

3	0	1	0	0	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

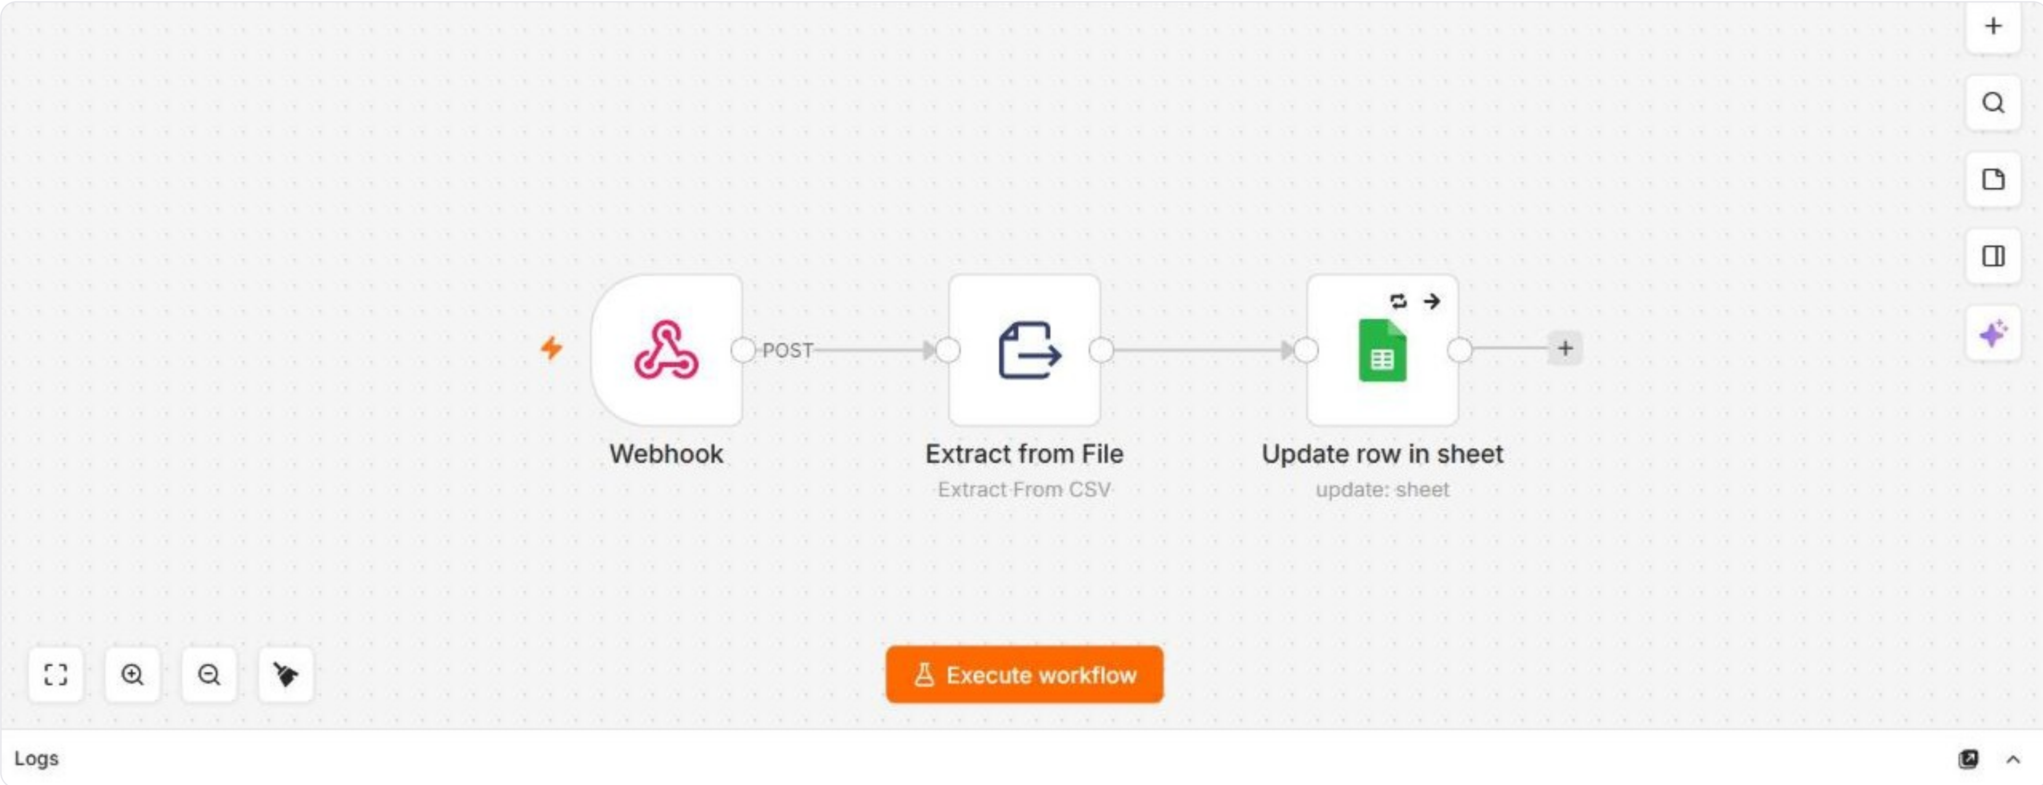

Wired to: Google Sheets

Audit database

The table everything else writes into. On a schedule it walks the pending rows, resolves which advertising campaign each vehicle belongs to with an LLM classification step, and keeps the sheet and the data table in sync.

16 Nodes	**2** AI nodes	**1** Triggers	**3** Branch points	**2** Code / mapping	**0** Sub-workflow calls

Wired to: Google Sheets · n8n DataTable · Claude (LLM gateway)

Form-to-sheet sync

A scheduled three-node sync that copies new form submissions from the data table into the sheet the operation reviews.

3 Nodes	**0** AI nodes	**1** Triggers	**0** Branch points	**0** Code / mapping	**0** Sub-workflow calls

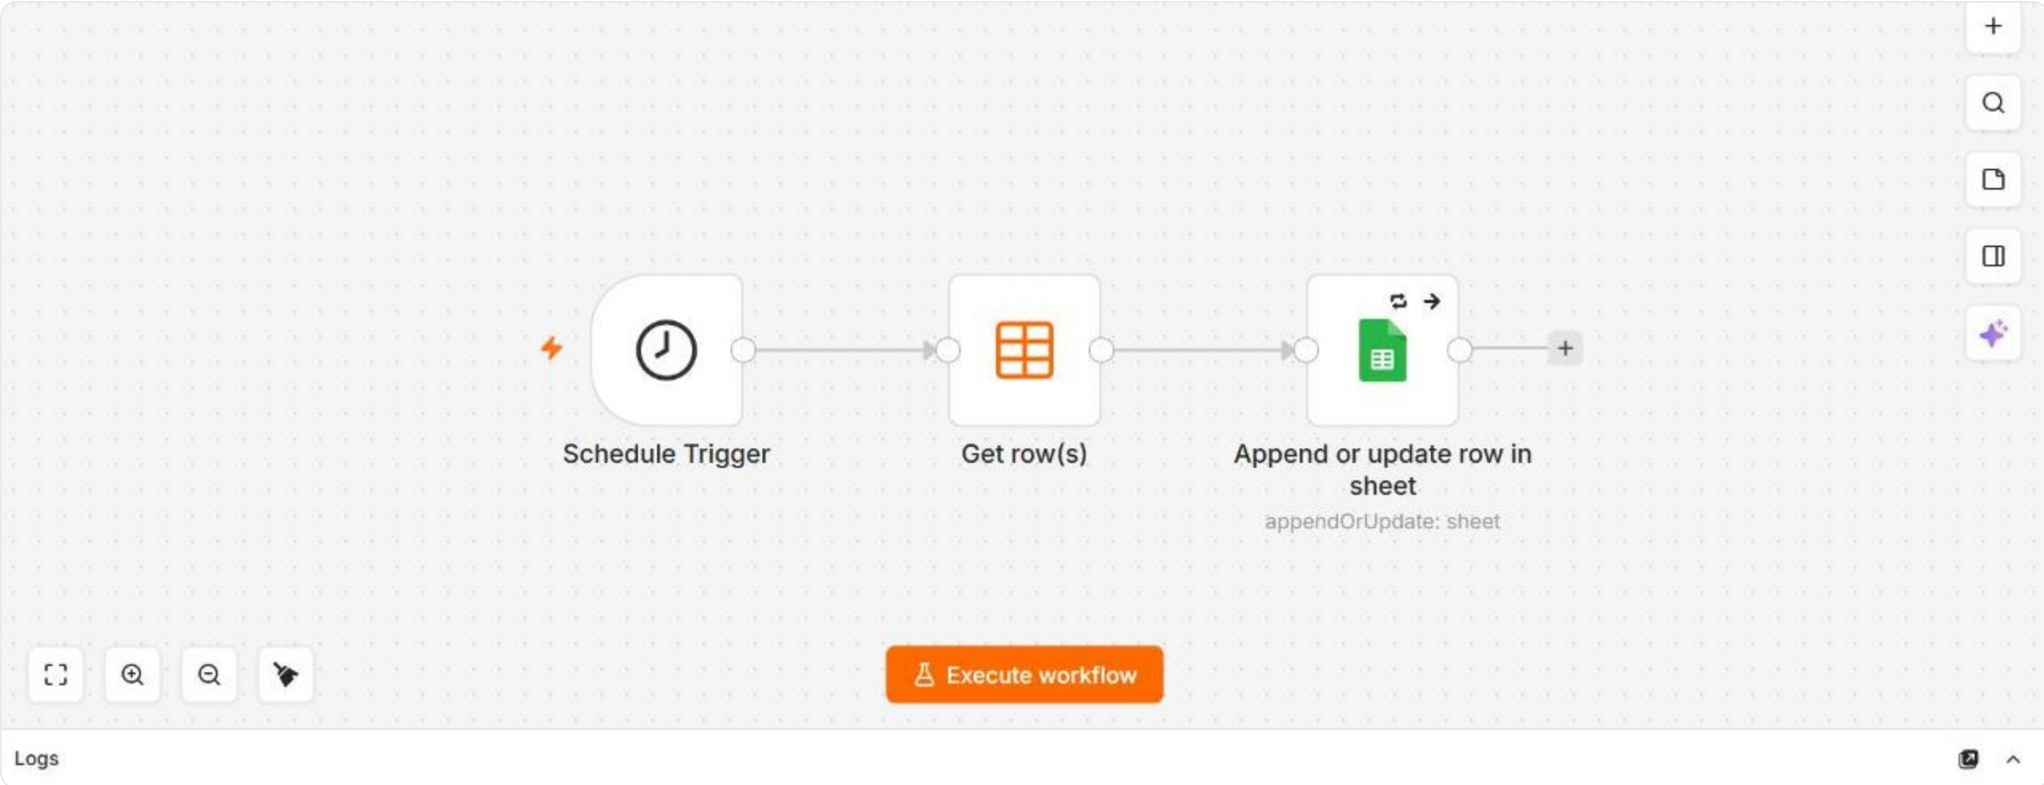

Wired to: n8n DataTable · Google Sheets

cabify

PROJECT

B2B — a corporate account, automated end to end

B2B Corporate Account · n8n

Two mail intakes, two AI agents that cross-check what the client asked against what the supplier committed to, and the flows that create the ride and pre-assign the named driver. Nothing is booked while the check disagrees.

4	**90**	**5**	**5**	**17**	**0**
Workflows	Nodes	AI nodes	Triggers	Branch points	Sub-workflow calls

Systems wired in: Claude (LLM gateway) · Google Drive · Google Sheets · HTTP / REST API · IMAP · Slack

Request intake

Watches the mailbox where the client sends its ride requests, pulls the spreadsheet attachment and drops it into Drive. This is what starts the rest of the chain.

4	0	1	0	1	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

Wired to: IMAP · Google Drive

Cross-checking AI agents

Two Drive triggers, two agents. One reads the client's request spreadsheet, the other reads the supplier's PDF report; both write into the same operational sheet, matching on the reference code, and post to Slack when a request has no report or the two disagree.

26	3	2	4	3	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

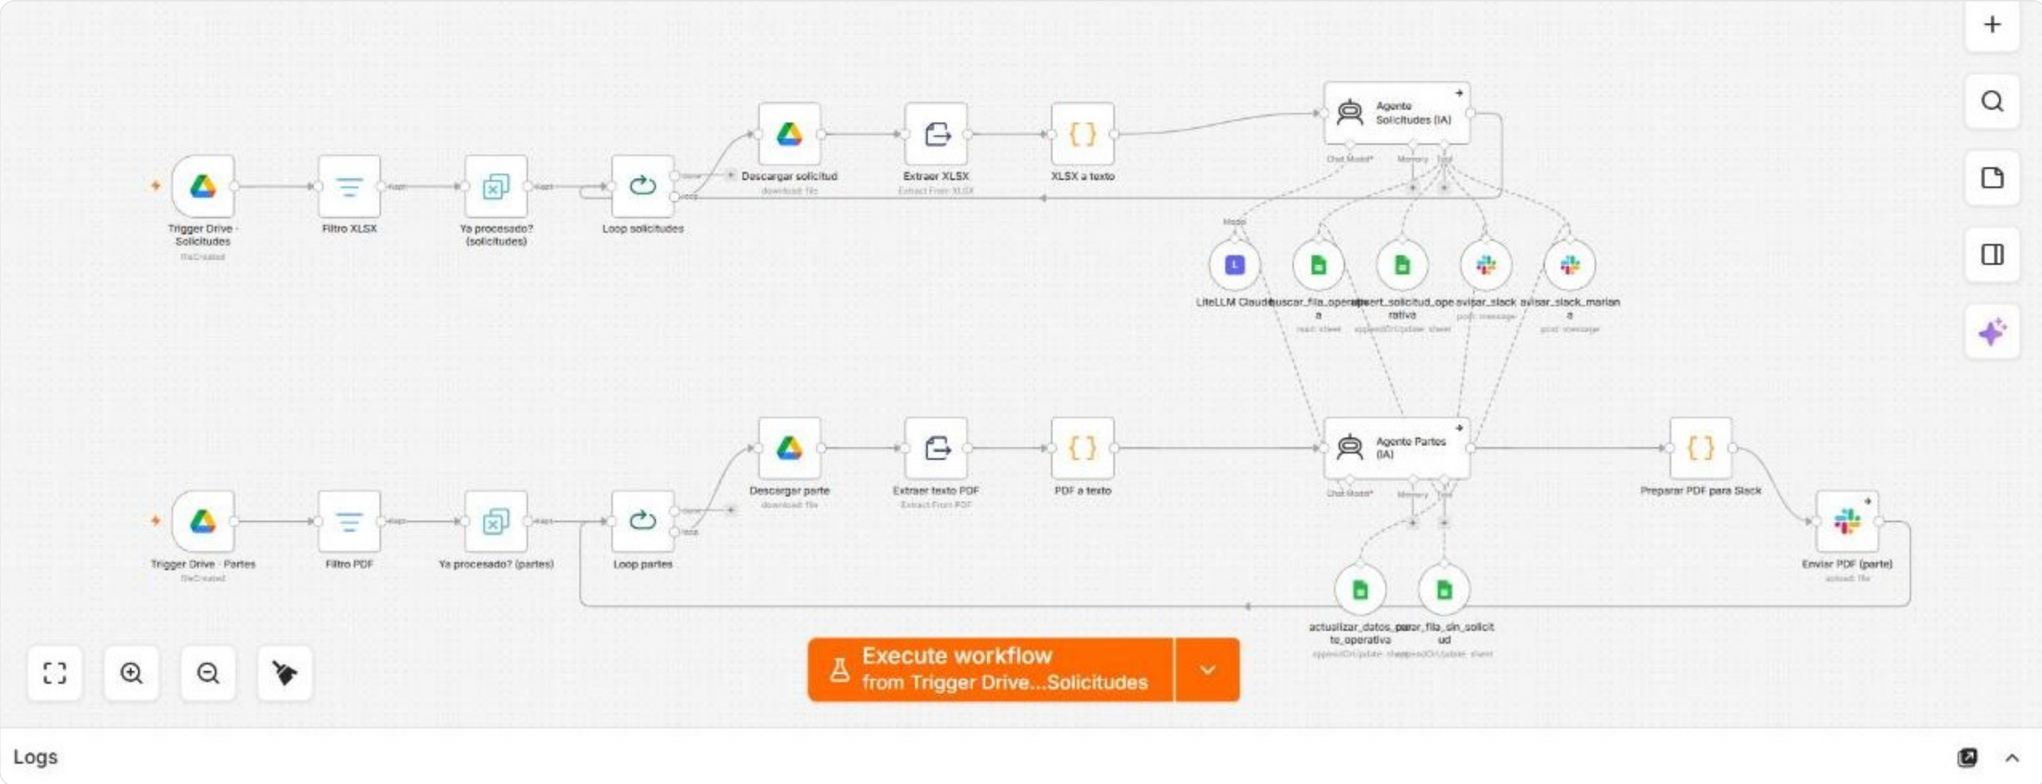

Wired to: Google Sheets · Slack · Google Drive · Claude (LLM gateway)

> **Why it is built this way.** The agents can only touch the sheet through four named tools — find row, upsert request, update report data, create orphan row. They cannot write arbitrary cells.

Ride creation and cancellation

Called by webhook once a row is validated: it authenticates, geocodes both ends, prices the trip, picks the contracted product, creates the ride and confirms it exists before writing the reference back and announcing it in Slack. The cancellation path mirrors it.

30 Nodes	**0** AI nodes	**1** Triggers	**6** Branch points	**6** Code / mapping	**0** Sub-workflow calls

Wired to: HTTP / REST API · Slack · Google Sheets

> **Why it is built this way.** Every write-back is verified against the API rather than assumed. If verification fails the row is marked as an error instead of silently looking booked.

Driver assignment

Assigns the specific driver the client asked for. It resolves the driver from the internal directory, or hands the candidate list to an AI agent when the name is ambiguous, then patches the ride and verifies the transfer before confirming in Slack.

30
Nodes

2
AI nodes

1
Triggers

7
Branch points

7
Code / mapping

0
Sub-workflow calls

Wired to: Slack · Google Sheets · HTTP / REST API · Claude (LLM gateway)

> **Why it is built this way.** The AI only chooses between candidates the directory already returned, and its choice is validated before the ride is touched.

cabify

PROJECT

COPs Control Board — the weekly control tower

COPs Control Board · n8n

One flow that reads the operational processes out of Tableau every week, has an LLM read each chart, maps the result into a master sheet and raises deviations in Slack by itself.

1	**108**	**10**	**2**	**1**	**0**
Workflows	Nodes	AI nodes	Triggers	Branch points	Sub-workflow calls

Systems wired in: Claude (LLM gateway) · Google Sheets · HTTP / REST API · IMAP · PDF parsing · Tableau

Weekly control board engine

One flow, roughly a dozen near-identical branches. Each signs in to Tableau, exports a view as an image, has an LLM read the chart and tag the numbers, flattens the result, maps it to the master schema and appends it. A separate branch does the same for the processes that arrive as email attachments.

108	10	2	1	52	0
Nodes	AI nodes	Triggers	Branch points	Code / mapping	Sub-workflow calls

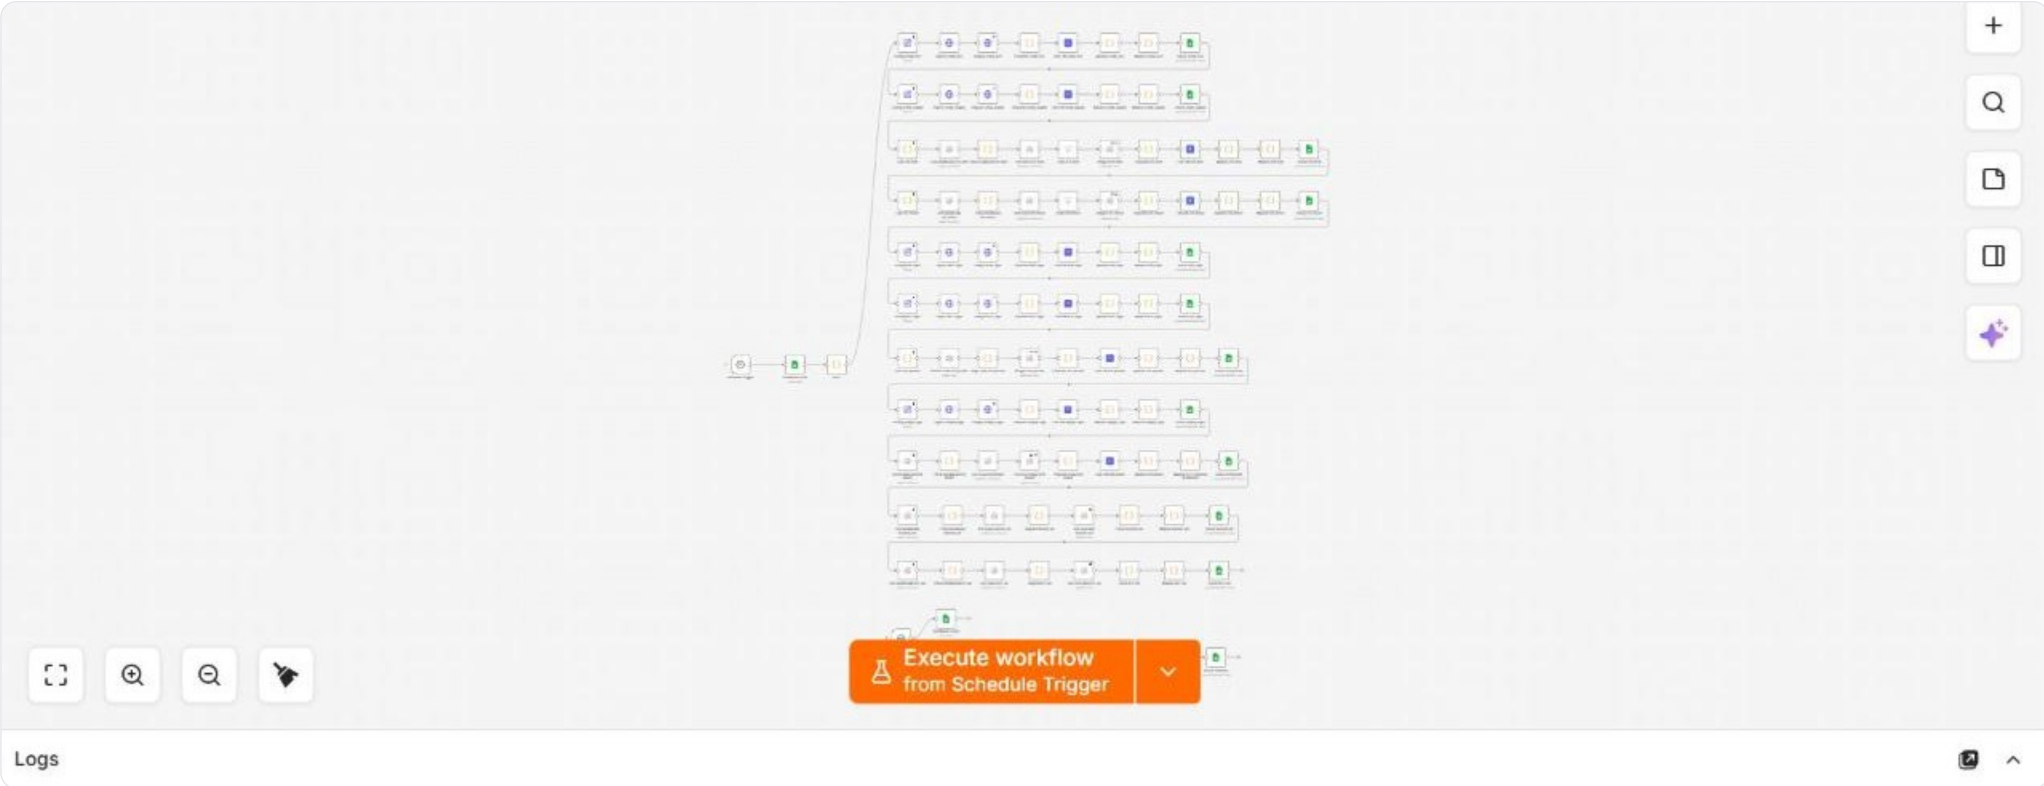

Wired to: Tableau · Google Sheets · HTTP / REST API · IMAP · PDF parsing · Claude (LLM gateway)

> **Why it is built this way.** The branches are identical by design. Adding the thirteenth process means copying a block and changing a config node, not rewriting the flow.

PROJECT

CabiBot — the public Help Center agent

CabiBot · Help Center AI Agent · n8n

Live support on the public Help Center: conversational memory, live article search, a ticket-lookup fallback, structured escalation with consent, and a WhatsApp escape hatch when a login simply cannot be unblocked.

1	13	4	1	0	2
Workflows	Nodes	AI nodes	Triggers	Branch points	Sub-workflow calls

Systems wired in: Claude (LLM gateway) · Google Docs · HTTP / REST API · n8n DataTable

Help Center AI agent

The public support agent. Redis memory holds the conversation, a search tool queries the Help Center article API live, and a ticket tool takes over when the articles cannot solve the case. Every exchange is logged to the usage table.

13 Nodes

4 AI nodes

1 Triggers

0 Branch points

2 Code / mapping

2 Sub-workflow calls

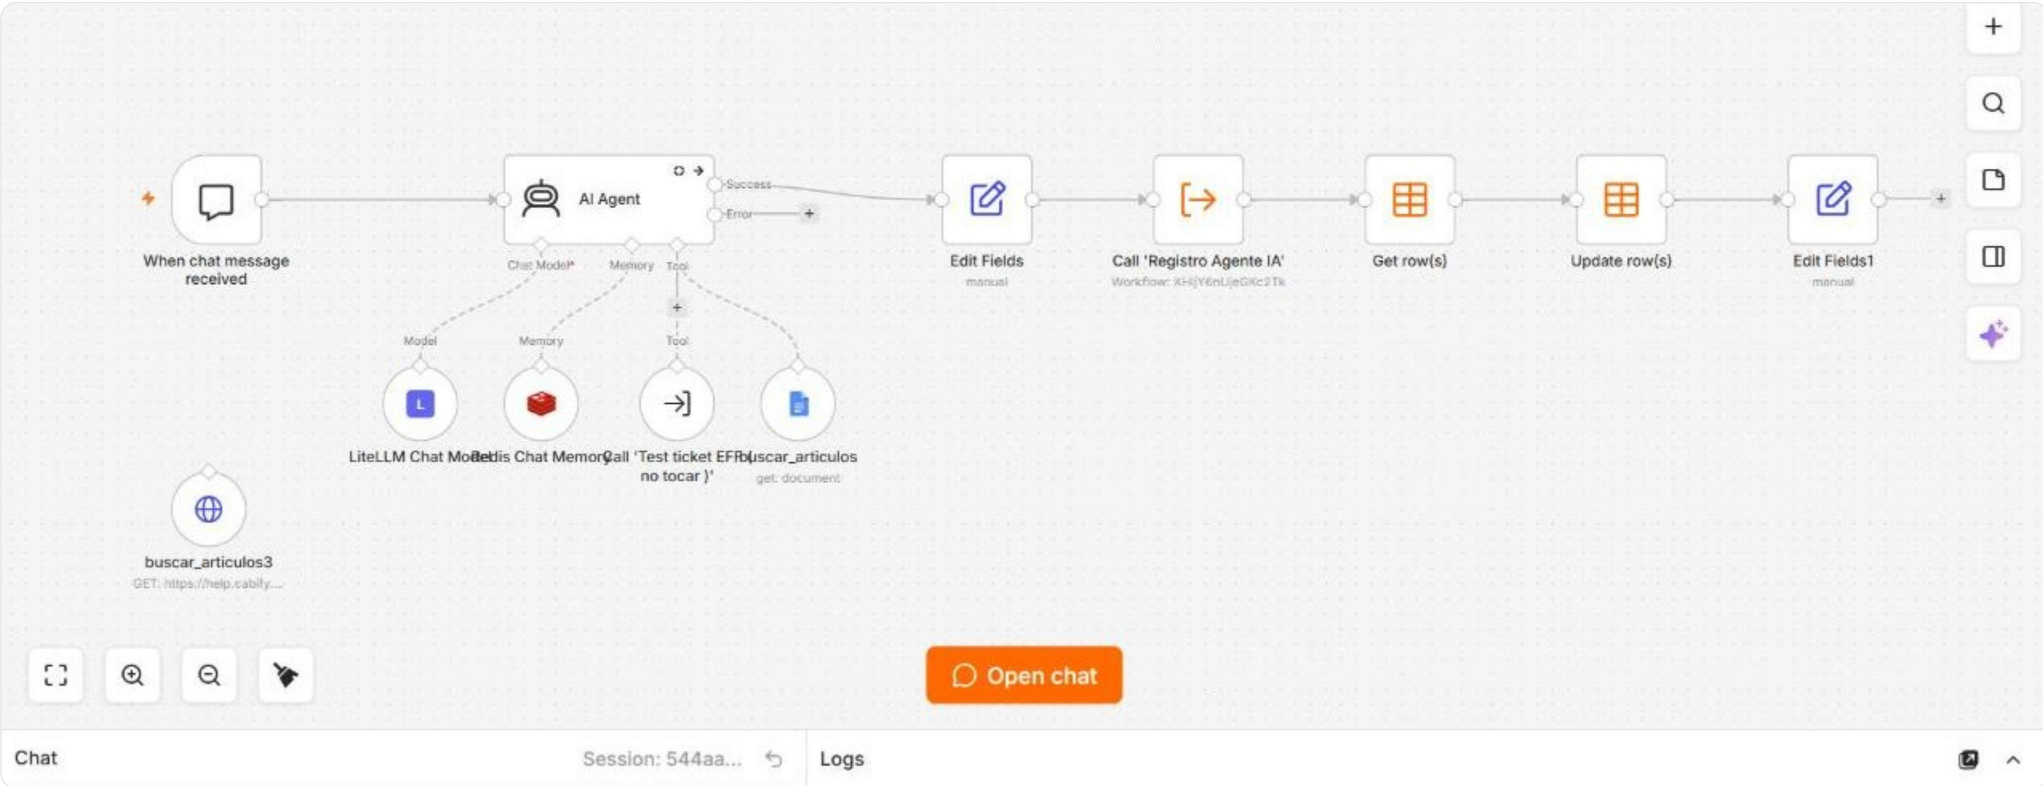

Wired to: n8n DataTable · HTTP / REST API · Google Docs · Claude (LLM gateway)

Why it is built this way. Article search runs against the live API, not a snapshot, so the agent cannot answer from documentation that has already changed.